

Doctor Strange

Primary protector of Earth
....adventures in bizarre worlds
and twisting dimensions



Cher

Goddess of Pop
....If I could find a way



David

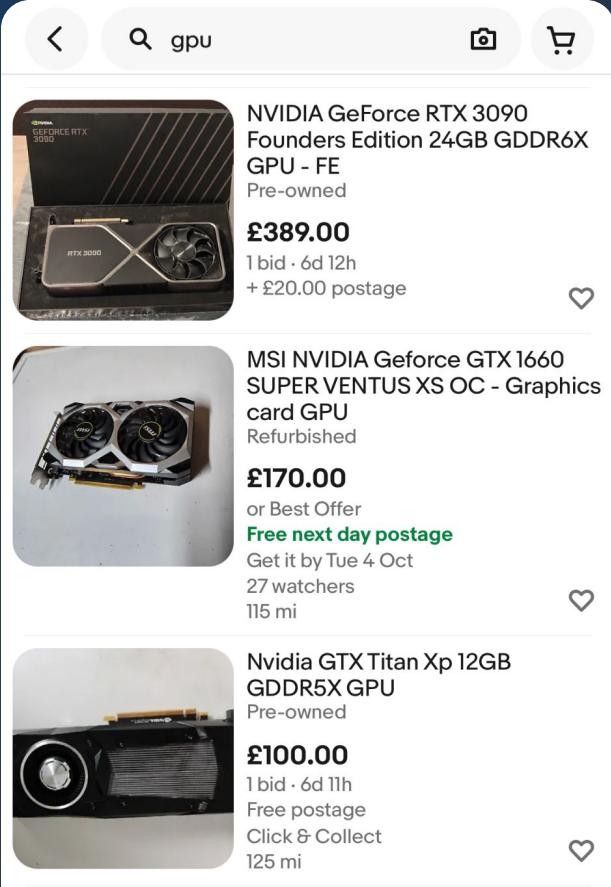
25 years in tech industry
....now professional Bug Hunter!



My house



Things I'm selling on eBay



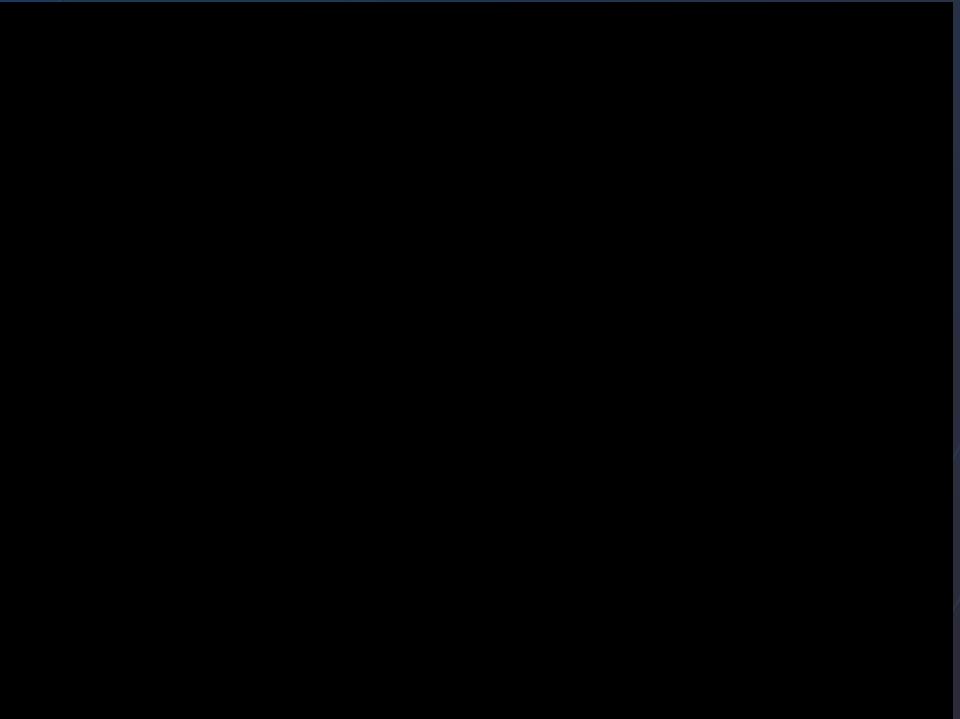
A screenshot of an eBay mobile search results page for "gpu". The search bar at the top shows the query "gpu". Below the search bar are three listing cards, each featuring a product image, title, price, and some descriptive text.

NVIDIA GeForce RTX 3090 Founders Edition 24GB GDDR6X GPU - FE
Pre-owned
£389.00
1 bid · 6d 12h
+ £20.00 postage

MSI NVIDIA Geforce GTX 1660 SUPER VENTUS XS OC - Graphics card GPU
Refurbished
£170.00
or Best Offer
Free next day postage
Get it by Tue 4 Oct
27 watchers
115 mi

Nvidia GTX Titan Xp 12GB GDDR5X GPU
Pre-owned
£100.00
1 bid · 6d 11h
Free postage
Click & Collect
125 mi

Or this?





Steven de Guzman

All my jobs include ensuring quality · Author has 1.5K answers and 7.1M answer views ·

In an imperfect IT world, software testing is difficult due to the **unfair expectation...**



Youji Hajime

Ex game tester, currently a compliance officer in completely unrelated field.
Author has 1.4K answers and 7.5M answer views · 6y

... Fundamentally, this is **impossible requirement** because “absence of evidence is not the evidence of absence (of error, bugs, breach, etc)”...



Werayut Nasawaeng

10++ years in Software Quality Domain · Author has 751 answers

... Software Testing is always never easy because we can't have 100% test coverage. We couldn't test **every cases possible**. And we will need to deal with time and resources constraints.



Sridevi Balla

6y

... The real issue is really **reporting bugs nicely**. Let me tell you a story where Wife is the Developer and the Husband is the QA...



Richard Dutcher

Former Enterprise Programmer at ADP (company) (2004–2013) · 3y

... So, testers have to think about **all possible scenarios** where issues may arise and ensure they are handled by the code.



Charitha Kankanamge

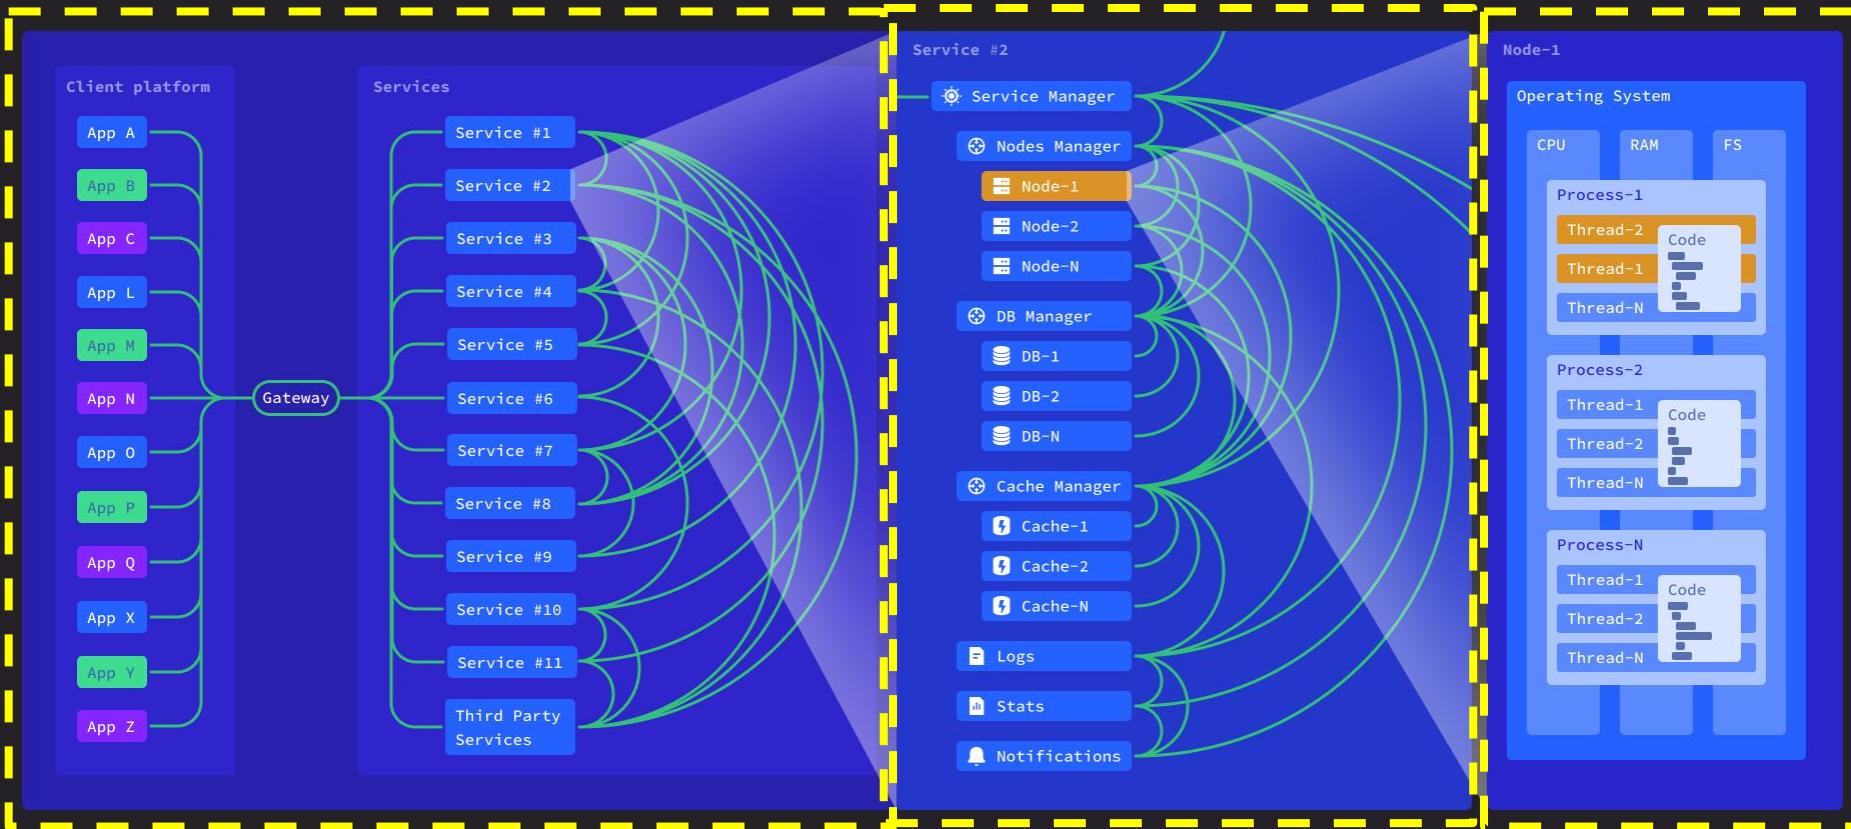
Manager, Software Quality at Amazon · 6y

... features which have been implemented without **considering testability** in advance...

Ethereum Testing

- Unit testing
- Testnets
 - Shadow-forks and testnet merges
 - Hive, Kurtosis, Antithesis
- Fuzzing
 - Coverage guided fuzzing on multiple portions of all CLs
 - Modified CL and EL clients for local-testnets and attack-nets.
- Static Analysis
 - manual code audits of clients, code queries (semgrep, codeql..)

Distributed Systems...



So what did we do...

- **Auto-generate** network, container, and other faults
- Fuzz **large systems**, not just the individual applications
- Ran all software **deterministically, across all clients at the same time**
- Used strategies to seek **rare events** (code coverage, events, log messages, etc.) and explored them as much as possible – better than random
- Used a **great exploration tool set**

Individually Tested

Consensus Clients



Lighthouse



Teku



Nimbus



Prysmatic Labs



Lodestar

Execution Clients



Go Ethereum



Nethermind



Hyperledger Besu



Erigon

Testnets & Shadow Forks

Consensus Clients



Lighthouse



Teku



Nimbus



Prysmatic Labs



Lodestar

Execution Clients



Go Ethereum



Nethermind



Hyperledger Besu



Erigon

One Simulation, many Parallel Universes



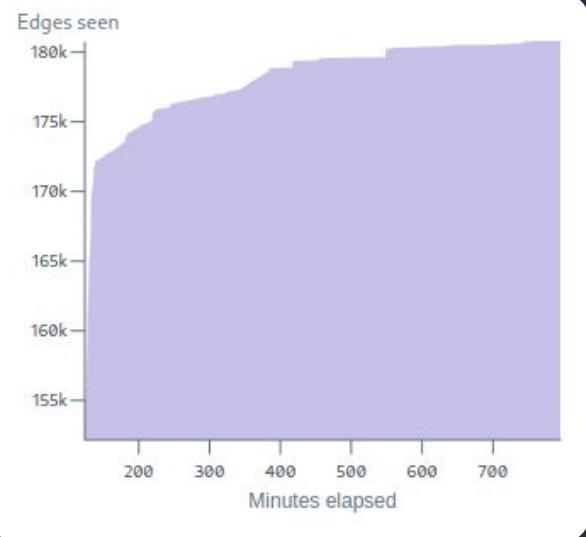


Test Branches
858

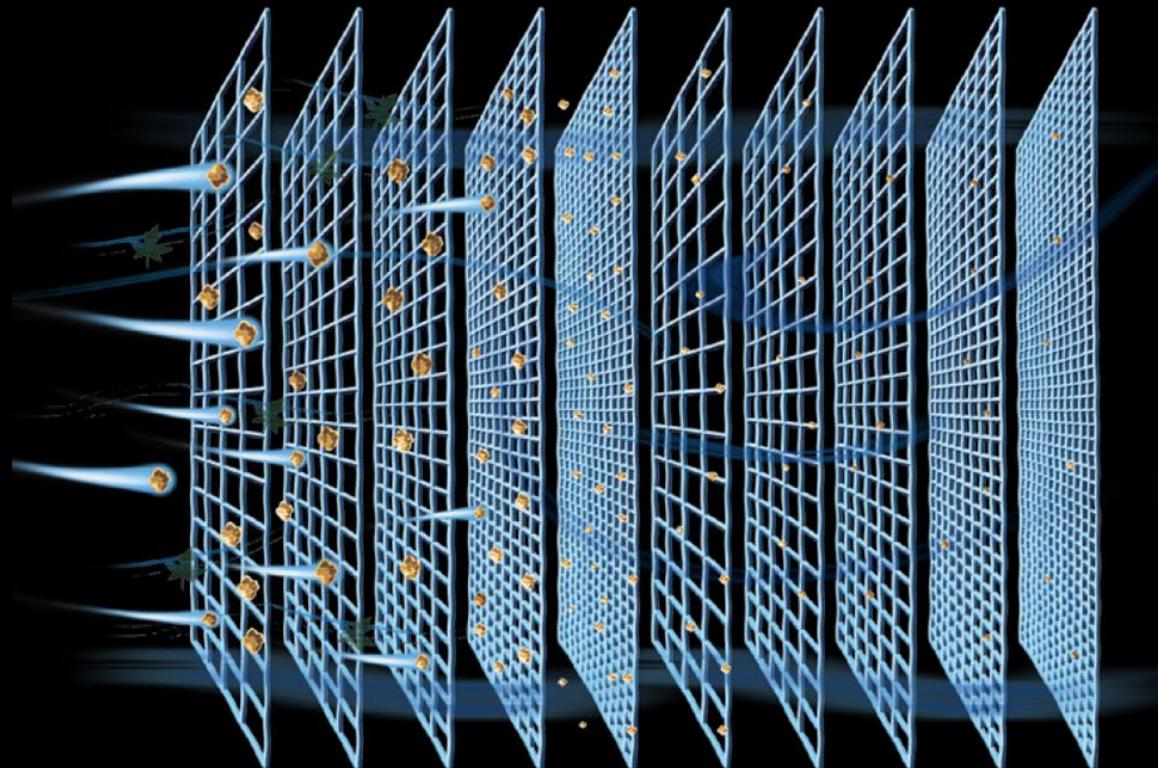
Test Hours (hrs)
536

Wall Clock (hrs)
13

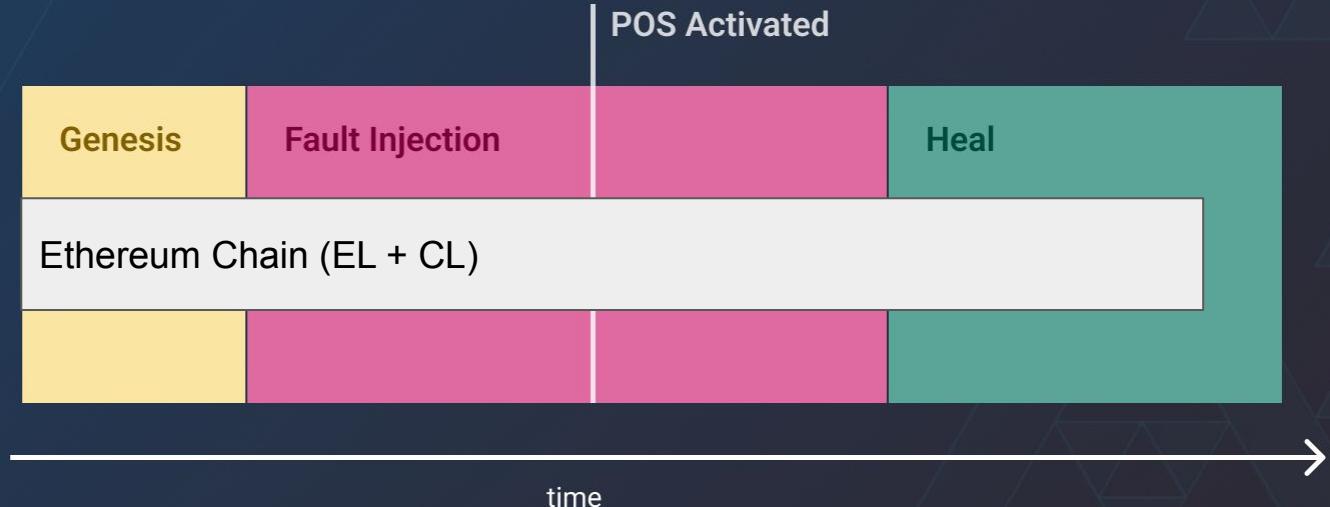
Edges Seen (unique)
180,802



Reducing the risk of exploits at every step



Simulating the Merge



Faults

Network: Partitions, Delays, Drops

Nodes: Stops, Pause, Kills

Threads: Pause, Release

31 years of 24x7 nonstop testing

Over 50million edges

45 validated errors

33 logged bugs



How about an example?

Report on test suite “Ethereum - Multi-client--kilnv2_inst Fault Tolerance Experiment”

Conducted on Wednesday, July 27th, 2022 at 5:23 PM

Nightly experiment expected to run for 16 hours with the latest images in the registry matching the given tags. Client images are expected to have been instrumented as the experiment is set up to use a coverage strategy for exploring the states.

► Experiment configuration details

Test Results

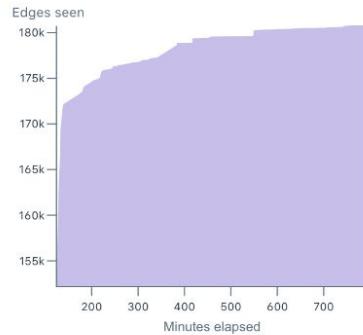
- Pass At some point workload started successfully
- Pass At some point there is at least one phase 3 success
- Pass At some point geth reports leaving PoW stage
- Pass Ultimately no panics
- Fail Ultimately no seg faults

Counterexamples

```
timestamp: 2402.57
source: service_nimbus-client-geth-0
rollout_id: 608
state_id: 119806
vtime: 10318950303345
text: SIGSEGV: Illegal storage access. (Attempt to read
from nil?)
```

- Pass Ultimately no impossible re-orgs
- Pass Ultimately no nimbus panics
- Pass Ultimately no lighthouse panics
- Pass Ultimately no prysm panics
- Pass Ultimately no lighthouse segfaults
- Pass Ultimately no lodestar segfaults
- Fail Ultimately no nimbus segfaults
- Pass Ultimately no prysm segfaults
- Pass Ultimately no teku segfaults

Test Branches	Test Hours (hrs)
858	536
Wall Clock (hrs)	Edges Seen (unique)
13	180,802




```
160 [service Prysm-client-nethermind-0] [I] time="2022-05-27 11:29:36" level=debug msg="gRPC request finished." backend=[] duration=93.338072ms method="/eth
161 [service Prysm-client-nethermind-0] [I] time="2022-05-27 11:29:36" level=warning msg="Fee recipient is currently using the burn address, you will not be able to withdraw funds from this account." backend=[] duration=117.162228ms method="/eth
162 [service Prysm-client-nethermind-0] [I] time="2022-05-27 11:29:36" level=debug msg="gRPC request finished." backend=[] duration=117.162228ms method="/eth
163 [service Prysm-client-nethermind-0] [I] time="2022-05-27 11:29:36" level=debug msg="gRPC request finished." backend=[] duration=22.352μs method="/ether
164 [service Prysm-client-nethermind-0] [I] time="2022-05-27 11:29:36" level=error msg="Could not request attestation to sign at slot" error="rpc error: connection closed" backend=[] duration=22.352μs method="/ether
165 [service Prysm-client-nethermind-0] [I] time="2022-05-27 11:29:36" level=debug msg="gRPC request finished." backend=[] duration=22.352μs method="/ether
[I] time="2022-05-27 11:29:36" level=error msg="Could not request sync message block root to the execution engine"
[I] panic: runtime error: invalid memory address or nil pointer dereference
[signal SIGSEGV: segmentation violation code=0x1 addr=0x0 pc=0x14a64f9]
[I]
[I] goroutine 266 [running]:
[I] github.com/prysmaticlabs/prysm/beacon-chain/blockchain.(*Service).notifyForkchoiceUpdate()
[I] /git/src/github.com/prysmaticlabs/prysm/beacon-chain/blockchain/execution_engine.go:193 +0x1519
[I] github.com/prysmaticlabs/prysm/beacon-chain/blockchain.(*Service).notifyEngineIfChangedHead()
180 [service Prysm-client-nethermind-0] [I] panic: runtime error: invalid memory address or nil pointer dereference
181 [service Prysm-client-nethermind-0] [I] [signal SIGSEGV: segmentation violation code=0x1 addr=0x0 pc=0x14a64f9]
182 [service Prysm-client-nethermind-0] [I]
183 [service Prysm-client-nethermind-0] [I] goroutine 266 [running]:
184 [service Prysm-client-nethermind-0] [I] github.com/prysmaticlabs/prysm/beacon-chain/blockchain.(*Service).notifyForkchoiceUpdate(0xc000342540, {0x21d8c
185 [service Prysm-client-nethermind-0] [I] /git/src/github.com/prysmaticlabs/prysm/beacon-chain/blockchain/execution_engine.go:193 +0x1519
186 [service Prysm-client-nethermind-0] [I] github.com/prysmaticlabs/prysm/beacon-chain/blockchain.(*Service).notifyEngineIfChangedHead(0xc000342540, {0x21
```

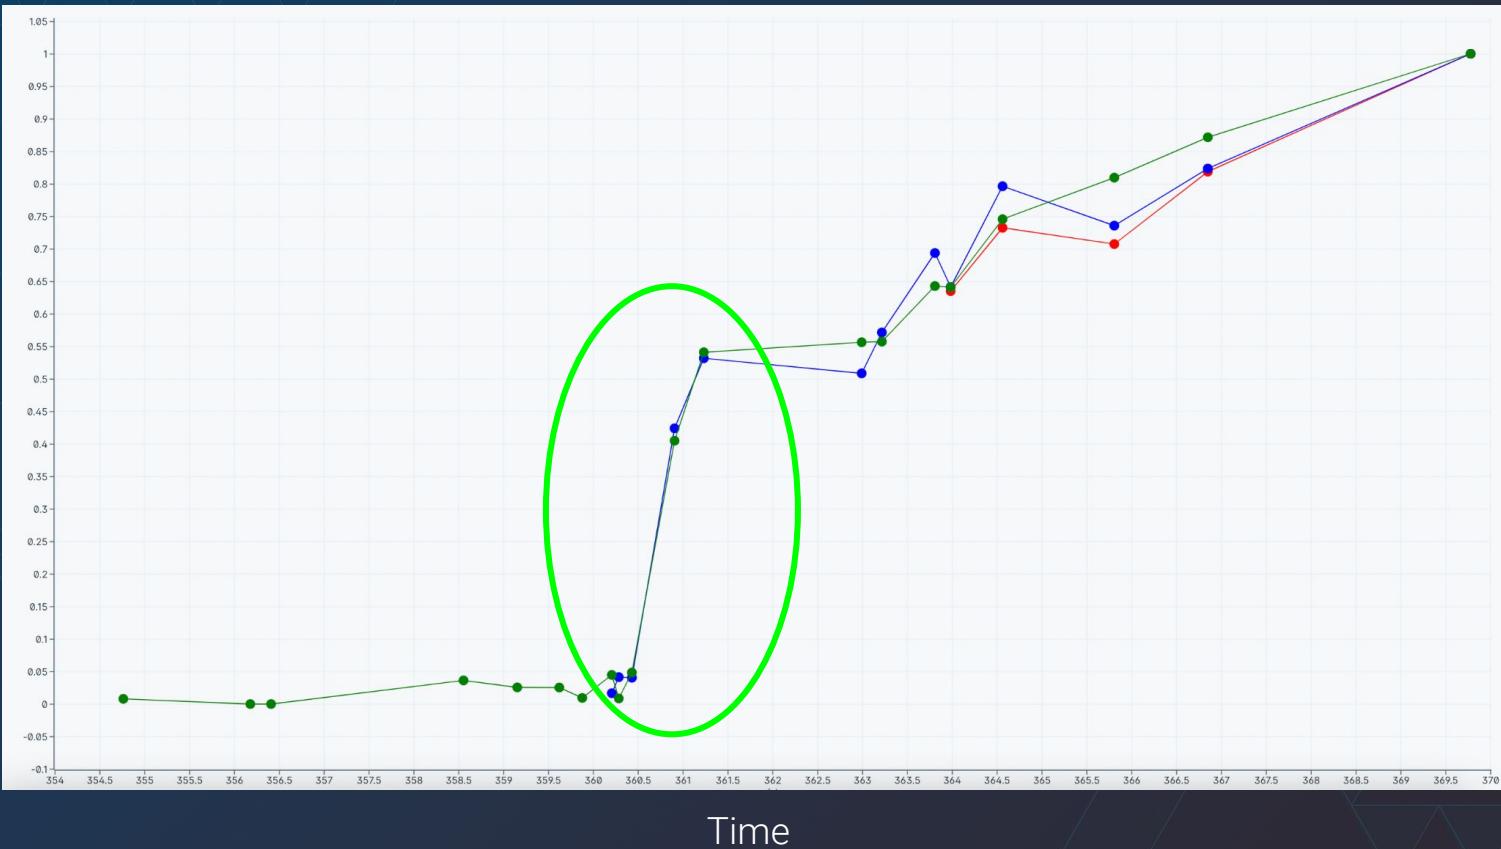
♪ If I could turn back time ♪



Look back n seconds ... and turn on packet
capture

But how far should you look back?

Probability of Bug



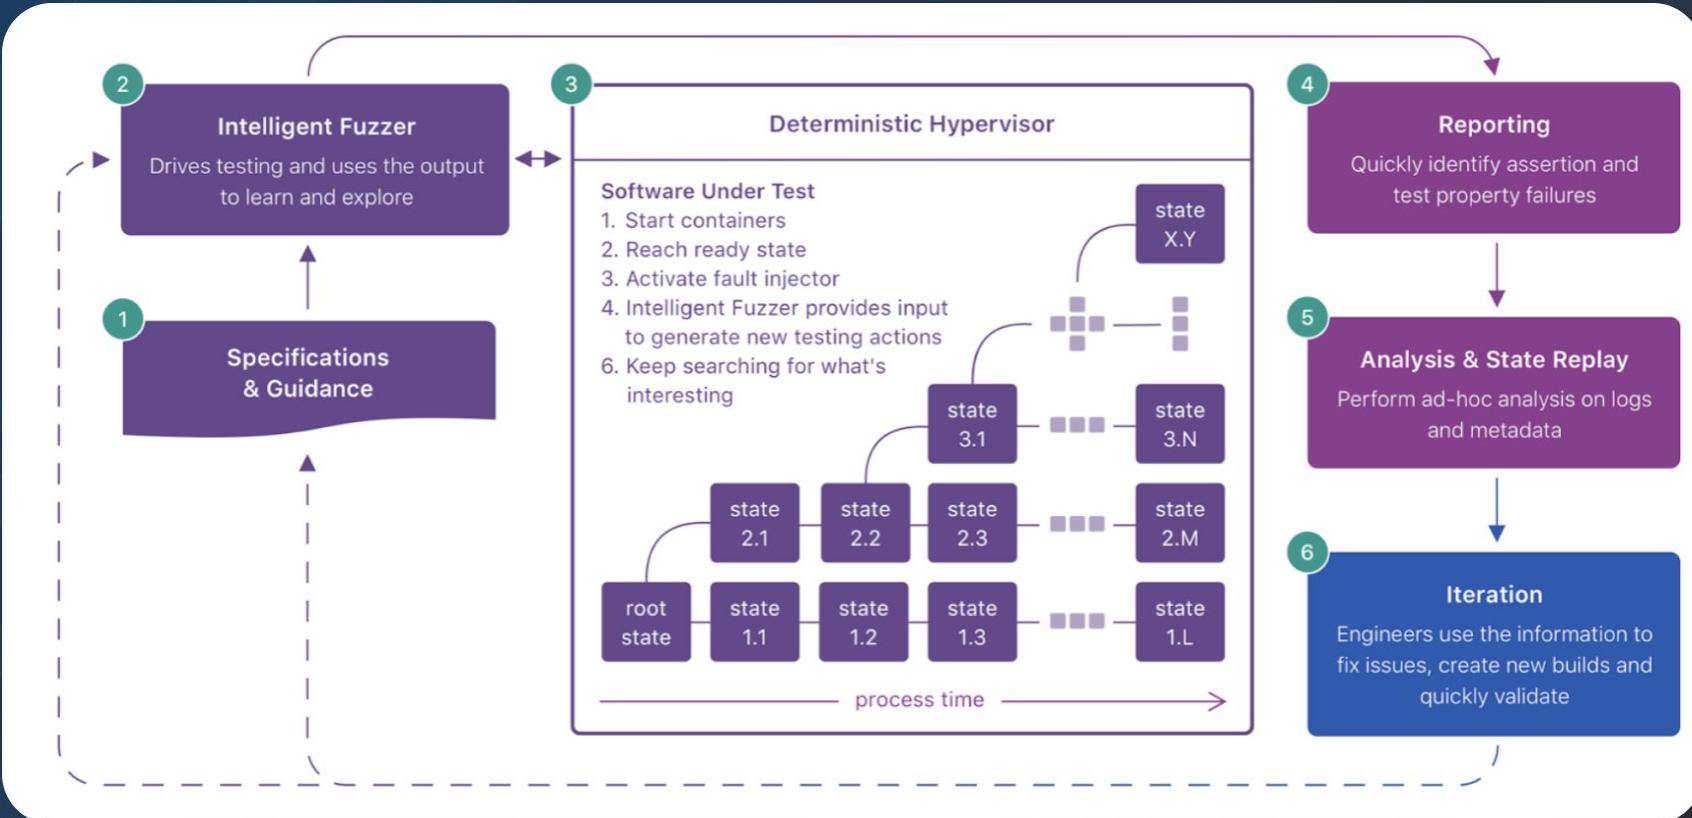
What's Next

- Post Merge Testing on Stable Branches
- EIP 4844 testing
 - Dank Sharding
 - Withdrawals
- Networking Testing Simulations (Gossip)
- API Fuzzing
- Byzantine / Malicious Client
- Merge code cleanup

Q & A

antithesis.com

tg: @signal1crypto



Reach vs Coverage

